

# Claude Multi-Account Fine Tuning

## llmctl Integration — Quick Start

1. Prerequisites
2. Quick start
3. The alias only exists while the model is actually running
4. Switching models
5. Two honest warnings you may see
6. Something not working?

## llmctl Integration — Quick Start

llmctl is a separate, sibling project that runs local LLMs on your own machine — each model profile as its own independent server, serving an OpenAI-compatible API on a fixed port. `claude-providers` automatically recognizes every `llmctl` profile you have **running right now** and exposes it as a normal, launchable Claude Code (and Kimi Code) alias — no manual provider-file editing, no separate registration step.

---

### 1. Prerequisites

| Need                                                         | Why                                                                                                                                                                    |
|--------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| llmctl installed, with at least one model profile downloaded | this integration discovers and uses llmctl — it does not install, configure, or manage llmctl itself (see llmctl's own docs for llmctl setup / llmctl models download) |
| The toolkit installed                                        | provides claude-providers, the alias file, the launch wrappers                                                                                                         |
| jq, curl, ss                                                 | catalog parse + liveness probe + bind-address check                                                                                                                    |

## 2. Quick start

```
llmctl start fast          # start a profile directly via
    llmctl
claude-providers sync      # re-sync claude_toolkit
source ~/.local/share/claude-multi-account/aliases.sh  # or open
    a new shell
claude-providers list | grep llmctl-  # confirm llmctl-fast
    appears
claude llmctl-fast -p "hi"  # launch Claude Code against it
```

The alias is always named `llmctl-<profile>` — never a bare profile name — so it is immediately distinguishable from a native account or any other provider alias. The same profile is also reachable through Kimi Code as `kimi-llmctl-<profile>` (e.g. `kimi-llmctl-fast`), using the exact same underlying detection.

## 3. The alias only exists while the model is actually running

`claude-providers sync` re-detects every `llmctl` profile on every run. If you stop the profile (`llmctl stop fast`), the next sync removes the alias — there is never a stale alias pointing at a model that is no longer there. If `llmctl` is not installed, or installed but serving nothing, `claude-providers sync` behaves exactly as if `llmctl` were absent: no error, no broken alias.

## 4. Switching models

Launching a **different** `llmctl-<profile>` alias than the one currently active switches the model automatically — the toolkit hands this off to `llmctl`'s own exclusive switch (stop the previous profile, start exactly the one you asked for). You do not need a separate switch command on the `claude_toolkit` side; just launch the alias you want:

```
claude llmctl-fast -p "hi"  # fast is now active
claude llmctl-coder -p "hi" # automatically switches: fast
    stops, coder starts
```

If the target profile does not fit your host's currently available resources, the switch is refused with a clear reason, and whatever was previously active is left exactly as it was.

## 5. Two honest warnings you may see

- **LAN-exposure warning** — the model is reachable from your local network, not only from this machine. `llmctl` binds to all

network interfaces by default; the alias still works, this is just telling you who else can reach it.

- **Context-size warning** — this profile's real context window is smaller than what a CLI agent's own prompt + tool-schema overhead typically needs, so a real turn may fail with a context-exceeded error. The alias still works for lighter use; this is an honest heads-up about a specific profile's capacity, not a toolkit defect.

## 6. Something not working?

See the full [user guide](#) for switching details, both warnings in depth, and troubleshooting — and the [FAQ](#) for quick answers to the most common "why isn't..." questions.