

Claude Multi-Account Fine Tuning

Provider Aliases — User Guide

1. Prerequisites
2. Quick start
3. Commands
 - Useful flags
4. How a provider alias is named and what it runs
 - 4.1 Model tier — strongest paid if you have credit, strongest free if you don't
5. Config dirs, plugins, and shared state
 - Cross-alias session visibility (v1.5.0+)
 - Multi-alias system (v1.6.0+; free-tier-first default since v1.26.0)
6. Overrides — `scripts/providers/overrides.json`
`scripts/providers/key-aliases.json`
7. Verification — what sync actually checks (v1.14.0+)
8. Known limitation — default session color
9. Troubleshooting
 - 9.1 Vendor-prefixed model ids (fixed in v1.26.7)
 - 9.2 The launch token guards and your plugin surface
10. Remote distributed deployment & heavy testing
11. Command quick-reference
12. Individual provider notes
 - Kimi Code — OAuth subscription (kimi-for-coding, kimi-k3, kimi-k2p7, kimi-for-coding-highspeed)
 - z.ai Coding Plan (zai-coding-plan)
 - Xiaomi MiMo (xiaomi)
 - OpenCode Zen (provider id opencode, alias opencode-zen)
 - Chutes (chutes)
 - Poe (poe)
 - HelixAgent (helixagent) — local single-GPU model
13. Releasing — the mandatory live pre-release gate

Provider Aliases — User Guide

claude-providers turns every LLM API key you keep in your keys file into its own Claude Code alias, pointed at that provider's strongest model. It reuses the multi-account toolkit's machinery (shared plugins/history, automatic `.bashrc/.zshrc` wiring) so a provider alias behaves like a normal `claudeN` account, just talking to a different model backend.

It is **fully dynamic**: the provider list, base URLs, transports, and model IDs are all derived at run time from your keys file + the [models.dev](#) catalog. Nothing is hardcoded. Re-run it any time to pick up new keys or refresh existing providers.

1. Prerequisites

Need	Why	Install
The toolkit installed	provides <code>lib.sh</code> , the alias file, <code>sync-state</code>	<code>bash scripts/install.sh</code>
A keys file	one <code>export NAME_API_KEY=...</code> per provider	<code>~/api_keys.sh</code> (default) or <code>pass --keys-file</code>
<code>jq</code> , <code>python3</code> , <code>curl</code>	catalog parse + resolve + fetch	system package manager
<code>claude-code-router</code> (<i>only for routed providers</i>)	translate Anthropic↔OpenAI/Gemini	<code>claude-ccr-build</code> — builds the bundled Go router from <code>submodules/claude-code-router</code> . Do not <code>npm install -g @musistudio/claude-code-router</code> ; that is a different binary the toolkit treats as a <i>doppelgänger</i> (it lacks the <code>ccr restart</code> subcommand every route-apply needs).
<code>submodules/LLMsVerifier</code> built (<i>optional</i>)	full "does this model work?" verification	see §7

Your keys file is the **only** place secrets live. `claude-providers` reads the variable *names* from it (without executing it) to discover providers, and the launch wrapper reads the *value* at the moment you start a session. No key is ever written into the repo, the alias file, or the per-provider env files.

2. Quick start

```
claude-providers sync          # discover + create/refresh all
                                provider aliases
source ~/.local/share/claude-multi-account/aliases.sh  # or open
                                a new shell
claude-providers list          # see what you got
deepseek                       # launch Claude Code on DeepSeek's
                                strongest model
```

3. Commands

Command	Does
<code>claude-providers sync</code> (<i>default</i>)	Discover every LLM key, resolve it via <code>models.dev</code> , verify it with live sentinel + tool-calling probes (see §7), and create/refresh one alias per provider. Idempotent — safe to run repeatedly.
<code>claude-providers list</code>	Table of verified provider aliases only (safe to launch): alias name, provider, transport, strong + fast model.
<code>claude-providers list-all</code>	Every installed provider alias, any status.
<code>claude-providers list-faulty</code>	Only aliases with an issue (failed / unverified / pending) — the filtered-out set.
<code>claude-providers show <id></code>	Print the resolved env file for one provider.
<code>claude-providers remove <id></code>	Remove the alias + env file; back up and remove the config dir.
<code>claude-providers add --from-key VAR --id PROVIDER</code>	Register a key→provider mapping (when a key's name doesn't match <code>models.dev</code>), then sync.

Useful flags

- `--keys-file PATH` — use a different keys file (default `~/api_keys.sh`).
- `--no-verify` — skip verification; create aliases regardless.
- `--offline` — don't fetch `models.dev`; use the local cache (errors if none).

- `--dry-run` — print what would change; write nothing.

4. How a provider alias is named and what it runs

- **Alias name** defaults to the `models.dev` provider id (`deepseek`, `groq`, `mistral`, ...). Want something shorter like `dseek`? Set it in `scripts/providers/overrides.json` (see §6).
- **Strong model** → Claude Code's main model (`ANTHROPIC_MODEL`). Chosen as the provider's most capable: reasoning-capable first, then newest, then largest context — **within the tier your account can actually pay for** (see §4.1).
- **Fast model** → Claude Code's background model (`ANTHROPIC_SMALL_FAST_MODEL`), chosen as the cheapest tool-call-capable model in that same tier.
- **Transport:**
 - `native` — provider speaks the Anthropic API directly (`models.dev` npm is `@ai-sdk/anthropic`). The alias runs `claude` with `ANTHROPIC_BASE_URL` / `ANTHROPIC_AUTH_TOKEN` / `ANTHROPIC_MODEL` set.
 - `router` — provider is OpenAI-compatible / Gemini. The alias launches through `claude-code-router` (`ccr default-claude-code`), which translates the protocol.

4.1 Model tier — strongest paid if you have credit, strongest free if you don't

Model selection is **credit-aware**. This is a hard rule applied to every provider alias:

Credit state of that provider's account	What the alias runs
Credit / purchased tokens available	the strongest paid model the provider serves that passes verification
No credit	the strongest free model (free tier / \$0 cost) that passes the same verification
Unknown — can't be determined	treated as <i>no credit</i> : the free choice

Why "unknown" falls back to free. The two possible mistakes cost very different amounts. Choosing a paid model on an unfunded key gives you a 402/403 at launch and a dead alias right when you wanted to work. Choosing a free model on a funded key only costs some capability, and the next `claude-providers sync` promotes it as soon as the credit signal becomes readable. So an inconclusive billing probe, a provider with no balance endpoint, or a stale offline catalog all resolve conservatively.

It's a floor on cost, not a cap on quality. Inside whichever tier applies, the pick is still the *strongest* model. The rule narrows the candidate set; it does not lower the bar within it. It applies to both the strong and the fast slot, and to every alias `sync --multi` generates.

Verification does not relax for free models. A free model still has to pass the sentinel probe and the tool-calling probe (§7) before its alias is activated. "Cheapest that works" never degrades into "cheapest".

Your pin always wins. A `strong_model` / `fast_model` entry in `scripts/providers/overrides.json` (§6) overrides the automatic choice completely. Use it when you know exactly what you're paying for on a given provider and want the tier logic to stay out of it.

The enforcement mechanism (in `providers_resolve.py`, with matching detection in the bundled `LLMsVerifier`) landed alongside this section. The table and the rules above are the **behavioural contract**; check the source for the current flag and field names.

5. Config dirs, plugins, and shared state

Each provider gets `~/.claude-prov-<id>`, which symlinks the same shared items as your accounts (`projects`, `todos`, `plugins`, `settings.json`, `CLAUDE.md`, ...). So **every plugin you have in `claude1..N` is available in every provider session**, and history/projects are unified. `sync` also force-enables the always-on set in shared settings: **superpowers, systematic-debugging, frontend-design, code-review**.

Provider dirs are deliberately **excluded from account auto-detection**, so they never get merged into your real accounts' auth/identity and never interfere with `claude-unify` or `claude-add-account`. Your existing `claudeN` aliases keep working exactly as before, and you can still create new accounts the same way.

Cross-alias session visibility (v1.5.0+)

Sessions created under **any** alias — `claude1`, `claude2`, `deepseek`, `opencode`, `xiaomi`, or any other — are visible from **every** other alias via `/resume`. This works automatically:

1. **Before launch:** `claude-sync-state pull` merges every account's and provider's `.claude.json` into the launching dir. This includes `lastSessionId` (what `/resume` uses), `allowedTools`, MCP config, and all other project settings.

2. **After exit:** `claude-sync-state` push merges the post-session state back out, so the next alias to launch picks up the new session.
3. **Session files** are in `~/.claude-shared/sessions/`, which every alias symlinks to — so the session data itself was always shared; the new work ensures `.claude.json` project metadata is also merged.

Example workflow:

```
# Start a project as deepseek
cd /path/to/my-project
deepseek          # bare launch: auto-creates the project's
                  session
# ... work for a while, then exit

# Continue the same project as opencode-zen
opencode-zen      # bare launch: AUTO-RESUMES the same project
                  session — no /resume needed
```

A **bare** alias launch (no arguments) now auto-resumes — or first-time creates — the one long-lived session keyed to the project root, so switching aliases continues the same conversation automatically (see §12 of `../Claude_Multi_Account_Fine_Tuning.md` and `SESSION_COLOR.md`). You only need `/resume` to switch to a **different** session than the project's own.

Performance: adds ~1-2 seconds per launch (jq merge across all dirs). Same overhead that `claudeN` aliases already have.

Note: provider dirs are still excluded from `cma_detect_accounts` (used by `claude-unify` and `claude-add-account`). Only `claude-sync-state` includes them, so the existing account management is unaffected.

Multi-alias system (v1.6.0+; free-tier-first default since v1.26.0)

`sync` creates **one alias per provider** using the 2 strongest models, and then (v1.26.0+, ATM-860) runs the **per-model multi-alias phase by default**: it verifies each provider's **FREE-tier models** with real completion probes and creates **multiple aliases** — `provider`, `provider2`, `provider3`, etc. — each with a different pair of strong + fast models. Free-vs-paid derives from real data only (models.dev cost rows, `:free` ids, self-hosted local endpoints); a model whose tier cannot be derived is treated as **paid** and skipped unprobed (fail-safe on spend), recorded honestly under `skipped_models` in `<provider>_verified.json`. Paid models are **never probed by default** — that costs real money and is an explicit opt-in.

```

# Standard (default): 1 alias per provider + per-model aliases for
#   FREE models
claude-providers sync

# Opt in to probing paid/unknown-tier models too (spends real
#   money)
claude-providers sync --include-paid          # or
                        CMA_SYNC_INCLUDE_PAID=1

# Only the per-model phase (no single-alias sync); free-tier by
#   default
claude-providers sync --multi

# Skip the per-model phase entirely (legacy single-alias-only
#   shape)
CMA_SYNC_MULTI=0 claude-providers sync

# Options
claude-providers sync --multi --max-aliases 10 --min-score 20

```

Alias generation is idempotent and deterministic: ranking/pairing tie-breaks on model id, so re-running sync yields byte-identical alias env files and the same provider/provider2/... naming. Generated provider aliases live in the ~/.claude-prov-* namespace and launch via `cma_run_provider` — a class the account dispatcher refuses, so they can never shadow or outrank a native claudeN account alias.

How verification works (v1.14.0+):

Each model is tested via live HTTP probes against the provider's chat endpoint: a sentinel probe (sends "Reply with exactly: VERIFY_OK" and **requires** the response to contain VERIFY_OK) and a tool-calling probe (requires the model to actually emit a tool call). Models are scored on 7 dimensions (0-100):

Dimension	Weight	What it checks
Existence + valid response	25pts	Model exists, returns real content
Tool calling	20pts	Supports function/tool_calls
Reasoning	15pts	Has chain-of-thought / reasoning_content
Context window	15pts	≥8K tokens (log scale)
Streaming	10pts	SSE streaming support
Latency	10pts	Response under 2s (full) or 5s (half)
Free tier	5pts	\$0 input cost

Tool calling is a hard gate, not just score. A model that answers chat but never calls tools is marked **not verified**, no matter how high its score — Claude Code is entirely tool-driven, so a tool-less model always breaks at runtime.

Anti-bluff detection prevents false positives:

- HTTP 200 with error body (JSON error wrapped in success)
- Empty response content
- **200 responses that lack the VERIFY_OK sentinel** (proxy fallback, silent model swap, canned text)
- Boilerplate "I'm unable to" responses
- Models that claim capability but don't deliver

Alias pairing:

- Models sorted by score, paired 2 per alias (strong + fast)
- Odd count: last model used for both positions in last alias
- Single model: used for both positions
- Default max: 5 aliases per provider (configurable)

Verification cache: results cached for 24h to avoid re-testing on every sync. Cache stored in `~/.local/share/claude-multi-account/providers/verification_cache.json`. The cache carries a schema version; results written by older verification logic (e.g. before tool calling was required) are ignored automatically.

6. Overrides — `scripts/providers/overrides.json`

Per-provider pins. Empty by default. Any field you set wins over the `models.dev`-derived value:

```
{
  "deepseek": {
    "alias": "dseek"
  }
}
```

This is how you (for example) give DeepSeek the short alias `dseek` — without touching any code.

A `strong_model` / `fast_model` pin here also **overrides the credit-aware tier choice** described in §4.1. If you pin a paid model on a provider whose credit state the toolkit reads as empty or unknown, that pin is honoured — the tier logic never second-guesses an explicit operator decision, so the launch failure that follows an unfunded key is then yours to expect.

Do not copy transport/base_url from an example. As of **v1.19.0** every *hosted* provider — DeepSeek and Xiaomi included — is pinned to router transport on its OpenAI-compatible endpoint, because both were verified working there and a single uniform path is far easier to debug. Every transport pin in overrides.json today is router: deepseek, xiaomi, opencode, opencode-go, chutes, kimi-for-coding, hyper. Two of them:

```
"deepseek": { "transport": "router", "base_url": "https://api.deepseek.com" }
"xiaomi":    { "transport": "router", "base_url": "https://api.xiaomimimo.com/v1" }
```

Pasting an older "transport": "native" / .../anthropic block would revert that fix. Always check the live values in scripts/providers/overrides.json before overriding transport or base_url.

The one shipped native pin left is the **local** helixagent-native (scripts/providers/helixagent-native.json), which is a separate provider record rather than an overrides.json entry — it is not a counter-example for a hosted provider.

scripts/providers/key-aliases.json

Maps a key variable name to a models.dev provider id, for the cases where the name differs (e.g. CODESTRAL_API_KEY → mistral, ZAI_API_KEY → zai). Use claude-providers add --from-key VAR --id PROVIDER to append entries.

7. Verification — what sync actually checks (v1.14.0+)

By default sync verifies every provider with **two live probes** against its chat endpoint, using the exact model the alias will run:

1. **Sentinel probe** — asks the model to reply with exactly VERIFY_OK and requires that token in the response. A 200 without it (or with an error object smuggled into the body) is a bluff: the endpoint answered *something*, not the requested model.
2. **Tool-calling probe** — requires the model to emit a real tool call (tool_calls / tool_use). Claude Code is tool-driven; a model that only chats is useless in practice.

Verdicts:

- verified — both probes passed. The alias is created and launchable.

- failed — a definitive problem: auth/billing rejection (401/402/403), model missing (404), missing sentinel, error-in-200, or no tool call. The alias is **not activated** (and the launch gate refuses to run it).
- unverified — the probes were inconclusive (rate-limit 429, 5xx, timeout, no network, no key). The alias is still created but the **launch gate** refuses to start it until a later sync reaches a verdict (override with `--force` — not recommended).

This replaced the old best-effort check (a bare `GET /models`), which proved only that the key was accepted — not that the model can actually run Claude Code. Both Anthropic-native (`/v1/messages`) and OpenAI-compatible (`/chat/completions`) endpoint shapes are probed in their native format.

For an additional authoritative layer ("can this model genuinely see and describe my code?") build the LLMsVerifier submodule:

```
git submodule update --init --depth 1 submodules/LLMsVerifier
cd submodules/LLMsVerifier && go build -o bin/model-
verification ./llm-verifier/cmd/model-verification
```

Once `submodules/LLMsVerifier/bin/model-verification` exists, sync prefers it as the existence check. sync also runs the two-round semantic code-visibility check (sentinel + independent judge) when it is available; billing/auth rejections there demote the alias, while transient judge/infra errors are an honest SKIP that never demotes. (Requires the Go toolchain; the build reads keys from your keys file.)

The live-TUI layer (`claude-providers verify <id> --deep, runner verify_superpowers_tui.sh`) additionally **proves which backend served the turn**. Router-transport aliases share one `ccr Router.default`, and an alias whose `base_url` is the gateway itself skips its own rewrite and would otherwise inherit the previous provider's route — a PASS that says nothing about the alias under test. Every evidence file now records `# ROUTE-INTENDED:` and `# ROUTE-RESOLVED:` (the latter read after the launch), and the leg fails with `# FAIL: route-mismatch` when they differ, or `# FAIL: route-unknown` when the route is unreadable. Native-transport aliases record an explicit `n/a` and are not route-checked.

Two further refusals close the remaining holes. `.Router.background` is checked alongside `.Router.default`, because Claude Code dispatches background sub-requests of the same turn through it — a mismatch there is `# FAIL: route-mismatch-background`. And because reading the config file back only proves what it says, the leg also demands a **restart receipt** bracketing the launch (a fresh gateway listening on line in `~/claude-code-router/<alias-id>/service.log`, or a changed `service.json` in that same directory): the wrapper's `ccr restart` runs under `|| true` and can legitimately be refused, which would leave the previous provider serving while

the file reads back correct. Without a receipt the leg fails closed with # FAIL: route-unproven. For router-transport aliases jq is a hard precondition — without it the route is unreadable and the leg refuses rather than skipping.

Every ccr path above is **per alias**: each alias owns ~/.claude-code-router/<alias-id>/ with its own config.json, service.json and service.log. There is no single global set to inspect.

Two further refinements to what the leg *prints*, both of which matter when you are triaging rather than releasing. A timeout (rc 124) whose transcript carries "terminal_reason":"aborted_streaming" is now marked # FAIL: stream-aborted (output_tokens=... duration_ms=... bound=...s) and states that it is **not** a trust/overwrite dialog — the old message guessed a dialog for every rc-124 and was wrong for all of them on the run that motivated the change. And a layer-3 failure now quotes the driver's own reason instead of asserting "cannot see code / bluffed": driver exit 1 also covers definitive provider rejections (401/402/403/404) of the model under test, so the message can now say non-200 status 402: "Insufficient balance for request." and send you to the account rather than to the model.

8. Known limitation — default session color

The original goal was for each provider session to default its / color to **purple**. Investigation of Claude Code 2.1.195 found that / color is **session-scoped and TUI-only — it is never persisted to disk**, and there is no settings key, hook, or environment variable to set a default color. So this cannot be automated with the current Claude Code.

Workaround: type /color purple at the start of a provider session. purple is a valid color. If a future Claude Code adds a persistable color setting, sync will seed it automatically.

9. Troubleshooting

Symptom	Cause / fix
provider 'X' needs claude-code-router	Routed provider, the bundled ccr gateway is not installed → claude-ccr-build (builds the vendored Go router; needs the Go toolchain).
resolved ccr ... is not the bundled claude-code-router /	The ccr in use is stale or is the npm @musistudio doppelgänger — it lacks the ccr restart subcommand, so restart is read as a profile name. → claude-ccr-build, and npm rm -g

Symptom	Cause / fix
Profile not found	@musistudio/claude-code-router if that is where the other one came from.
\$NAME_API_KEY is empty	The key isn't set in your keys file.
offline and no models.dev cache	Run <code>claude-providers sync</code> once online to populate the cache.
A key didn't become an alias	It's unmapped (no models.dev match) or classified vcs/infra. Map it with <code>add</code> , or check <code>scripts/providers/evidence/mapping-report.md</code> .
Alias missing after sync	<code>source ~/.local/share/claude-multi-account/aliases.sh</code> or open a new shell. The catalog row for the model is self-inconsistent (<code>limit.output >= limit.context</code>). <code>derive_limits()</code> in <code>scripts/providers_resolve.py</code> always carves the output cap out of the context instead of trusting the published one, so such a row resolves to the same cap an absent output would; re-run <code>claude-providers sync</code> to rewrite the alias <code>.env</code> . A model the catalog has never heard of (an operator pin) no longer emits <i>no</i> limits either — an unknown context resolves to a conservative measured fallback (the catalog p10, narrowed by the provider's own p10, capped by that provider's widest window and floored so it stays usable), so both guards
400 "exceeded model token limit"	always exist. A <code>:free</code> row whose output budget exceeds its paid sibling's is only adjudicated at all when its output < context (where they are equal the row is just the common output-copied-from-context mislabel, which the carve already fixes), and its context is only lowered when at least three <i>distinct</i> providers other than the accused one , publishing the same model under a compatible vendor prefix, put it lower — then only to the lower median of what they publish. A lone suspicious row is left alone, because a wrongly-shrunk window is its own dead alias. Note the deliberate blind spot in the other direction: a genuinely throttled <code>:free</code> tier is corroborated at the <i>paid</i> value and left alone, since every peer record describes the paid tier.
# FAIL: route-mismatch in a --deep evidence file	The verification turn was served by a different backend than the alias under test (see §7). Re-run that leg on its own rather than after another router alias.

Symptom	Cause / fix
# FAIL: route-mismatch-background	Same, for the <code>.Router.background</code> entry: background sub-requests of that turn went to another backend (see §7).
# FAIL: route-unproven	The route was written but nothing proves the live gateway loaded it — the ccr restart was refused or never happened. Most often an authenticated gateway restarted without <code>CCR_API_KEYS</code> visible; re-run the leg with the gateway's keys in the environment.
Alias is verified, config looks right, every launch dies as an apparent timeout/hang	The vendor-prefix routing class fixed in v1.26.7 — see §9.1. Confirm with <code>tail ~/.claude-code-router/<alias-id>/service.log</code> : a burst of 503s at 1-4 ms each means the <i>local</i> gateway rejected the route before any network call, so no <code>--timeout</code> value helps. Rebuild the router (<code>claude-ccr-build</code>) if you are on an older build.
# FAIL: stream-aborted in a <code>--deep</code> evidence file	The bound expired <i>and</i> the CLI wrote a result with <code>terminal_reason=aborted_streaming</code> (the marker carries <code>output_tokens</code> , <code>duration_ms</code> , <code>bound</code>). Explicitly not a trust/overwrite dialog. Where it died is not determined by the leg — read the alias's <code>service.log</code> as above.
400 naming ... of tool input ... and a total over the window	Your enabled-plugin surface is larger than <code>CMA_TOOL_TOKEN_BUDGET</code> (default 80000), which inflates the derived output cap. See §9.2 — raise the budget, cut plugins, or move to a wider-context provider.

9.1 Vendor-prefixed model ids (fixed in v1.26.7)

A route is `provider,model` on disk, but the router also accepts a **client-supplied** selector in the request's `model` field, and it recognises *two* separators — comma and slash, leftmost wins (`submodules/claude-code-router/internal/router/selector.go`).

Once `lib.sh` began exporting `ANTHROPIC_MODEL` for the **router** transport too, Claude Code put the raw catalog id in the request body — `deepseek-ai/DeepSeek-V3.2-TEE`, `nvidia/nemotron-...`, `Qwen/...` — the router read the **vendor** as a provider name, found none configured in the per-alias config, and returned an unknown-provider error instead of falling through to `Router.default`. The gateway answered **503 in 1-4 ms — before any network call** — and retried with backoff until the launch bound expired. **16 of 24 verified aliases** were affected.

What made it hard to see: the on-disk route was correct the whole time ("kilo,nvidia/nemotron-..." parses fine, its comma comes first), so the config looked healthy while every launch died. And layers 1-3 of verification curl the provider's own base_url directly and never touch the ccr route path, so they stayed green throughout. **Only the layer-4 live launch exercises the real chain.**

The fix decides the slash form from **evidence, not from the separator**:

- if some configured provider actually serves the whole string, it is a catalog id and routes normally;
- if a vendor prefix *collides* with a real provider name (provider nvidia serving nvidia/nemotron-...), the whole id wins — the split half is absent from Models while the full string is present;
- if nobody serves it, the unknown-provider error still stands, so Ghost/whatever cannot silently reach Router.default.

Comma selectors are unchanged and still fail loudly. A comma is this project's own route syntax and appears in no model id, so an unknown provider really is a caller error; silently redirecting it to Router.default would bill an upstream you never chose.

9.2 The launch token guards and your plugin surface

The launch wrapper exports two co-derived guards for every alias with a known context: CLAUDE_CODE_MAX_OUTPUT_TOKENS (output cap) and CLAUDE_CODE_AUTO_COMPACT_WINDOW (compaction trigger). The window must reserve for the **tool payload** as well as text and output, because the endpoint counts all three against the same window:

```
window = context - output - CMA_TOOL_TOKEN_BUDGET          # budget
default 80000
```

and when that window would fall below CMA_MIN_COMPACT_WINDOW (default 120000), the wrapper buys the window back by lowering the output cap instead:

```
CLAUDE_CODE_MAX_OUTPUT_TOKENS = context - CMA_MIN_COMPACT_WINDOW -
CMA_TOOL_TOKEN_BUDGET
```

Be exact about which direction is dangerous. An *underestimate* does not merely compact earlier — because the output cap is derived by subtracting the budget, too small a budget **inflates** the cap and the launch dies with a 400 naming the real numbers. Measured live on a host with **227 enabled plugins**, real tool input **158112** against the 80000 default gave:

46536 text + 158112 tools + 62144 output = 266792 against a
262144 window

Critically, $158112 + 120000 > 262144$: on such a host **no output cap makes a 262144-context provider fit**, because the tool payload plus the minimum compaction window already exceed the whole context. The only remedies are reducing the enabled-plugin surface or moving to a wider-context provider.

Override per host with `CMA_TOOL_TOKEN_BUDGET`, either as a line in the provider's resolved `.env` or in the shell. To measure your own tool input, read it out of the 400 itself — the provider's message names it, e.g. you requested about 371727 tokens (199347 of text input, 70236 of tool input, 102144 in the output).

Verify the context before calling an alias "wide enough."
Read the resolved env file rather than any table in this guide, since the value is derived per selected model at sync time:

```
grep CMA_PROVIDER_CONTEXT_LIMIT ~/.local/share/claude-multi-  
account/providers/<alias>.env
```

A context too small to reserve the budget at all yields no window: that provider cannot serve Claude Code's tool prefix regardless, and the launch fails loudly rather than being lied to by a guard that cannot hold.

10. Remote distributed deployment & heavy testing

For heavy testing that depends on **real production services**, the toolkit can boot the full LLMsVerifier System on a remote host (registered in `config/containers/nezha.env`) via the `containers` submodule + `podman`, and run real per-provider verification through it.

- **Deployment guide + boot procedure:** `config/containers/llmsverifier/README.md` (services, ports, the exact `podman-compose` boot steps, and every root-caused deployment fix).
- **What runs:** the `llm-verifier` API + an observability tier (`prometheus` + `grafana` with an auto-provisioned datasource & dashboard + `node-exporter`).
- **Heavy testing:** the bundled model-verification tool runs a real "Do you see my code?" check against each provider's live API:

```
ssh <user>@<host> 'podman run --rm --env-file ~/helix-system/  
llmsverifier/.env \  
llm-verifier-mv:nezha --provider deepseek --model deepseek-  
chat --verbose'
```

- **Evidence:** every verification + deployment claim is captured under docs/qa/20260616-infra/ (proofs, sweeps, security, observability).

11. Command quick-reference

```
# Provider aliases
claude-providers sync                # discover keys -> create/
    refresh aliases
claude-providers sync --no-verify    # skip verification (faster)
claude-providers list                # alias | provider |
    transport | models
claude-providers show <id>           # one provider's resolved
    env
claude-providers remove <id>         # remove alias + env, back
    up config dir
claude-providers add --from-key VAR --id PROVIDER # register
    mapping + sync

# Launch a provider session (after: source the alias file or open
    a new shell)
deepseek                             # router provider -> claude
    via ccr
<router-provider>                    # routed provider -> claude
    via ccr
deepseek -p "your prompt"            # non-interactive print mode

# Accounts (unchanged, still works)
claude-add-account --alias claudeN   # add a Claude account
claude-list-accounts                 # status of all accounts
claude-unify                         # re-merge shared state

# Docs
claude-export-docs                   # regenerate HTML/PDF/DOCX

# Release gate (run before tagging a release)
claude-release-gate                  # sandbox suite + LIVE real-
    alias smoke
claude-release-gate --verify-providers # also run the full
    LLMsVerifier scan
```


12. Individual provider notes

Kimi Code — OAuth subscription (kimi-for-coding, kimi-k3, kimi-k2p7, kimi-for-coding-highspeed)

If you are signed into **Kimi Code** (the kimi CLI), every model your subscription serves becomes an alias automatically — **Kimi 3** (kimi-k3, 1M context, reasoning), **Kimi 2.7** (kimi-k2p7), kimi-for-coding-highspeed, and the account default kimi-for-coding. No API key is required; the OAuth session in ~/.kimi-code/credentials/kimi-code.json is used.

- **Discovery:** claude-providers sync queries GET /coding/v1/models with your OAuth token and emits one alias per served model (unioned with the models.dev catalog, because the listing under-reports). Each alias is verified with the strict sentinel + tool-calling + semantic pipeline.
- **Token freshness:** the OAuth token lives ~15 minutes. At every launch the wrapper reads the **live** credentials file (with expiry), refreshes via `kimi -p hi` when expired, and only then falls back to the sync-time snapshot — launches never die of a stale token.
- **kimi_proxy:** k3 enforces a "moonshot-flavored" JSON schema for tools (every \$ref must start with #/\$defs/). Claude Code's tool schemas would 400 without it, so all kimi-* launches route through a local normalizing proxy (scripts/proxy/kimi_proxy.py, installed by install.sh).
- **Precedence:** an OAuth subscription wins over KIMI_API_KEY / ApiKey_Kimi records for kimi-for-coding; the API keys remain the fallback on hosts without the OAuth session.

z.ai Coding Plan (zai-coding-plan)

The z.ai Coding Max-Yearly Plan provides access to Zhipu AI's flagship GLM models through a dedicated coding-optimized API endpoint. This is separate from the base z.ai plan — the Coding Plan endpoint uses `api.z.ai/api/coding/paas/v4` and includes **Free access** to glm-5.2 (the flagship with 1M context and reasoning) and glm-4.7 (the fast model with 204k context, reasoning, and tool_call support), among others.

How the alias works

The key variable ZAI_API_KEY in your keys file is mapped to the zai-coding-plan provider ID via scripts/providers/key-aliases.json. An override in scripts/providers/overrides.json pins the strong model to glm-5.2 and the fast model to glm-4.7, which are the optimal choices for coding workloads on this plan.

Transport is **router** — the alias launches through `claude-code-router` (`ccr default-claude-code`), which translates the Anthropic protocol to the OpenAI-compatible `z.ai` API.

Models available on the Coding Plan

All models are free on the Coding Max-Yearly Plan:

Model	Context	Reasoning	Tool Call	Notes
glm-5.2	1M tokens	Yes	Yes	Flagship — newest, most capable
glm-5.1	—	Yes	Yes	Intermediate
glm-5-turbo	—	—	—	Turbo variant of glm-5
glm-5	—	—	—	Base generation 5
glm-4.7	204k tokens	Yes	Yes	Fast model — reasoning + tool_call
glm-4.6	—	—	—	Mid-range
glm-4.5-air	—	—	—	Lightweight
glm-4.5	—	—	—	Base model

Note on glm-5.2 rate limits: The Coding Max-Yearly Plan enforces a Fair Usage Policy on the flagship model: rapid-fire requests (e.g., repeated programmatic calls at short intervals) may trigger a temporary rate limit (error code 1313). Normal interactive use via `zai-coding-plan` at human pace is unaffected.

Setup

Everything is already configured on this host, but for a fresh install:

1. Ensure `ZAI_API_KEY` is exported in your keys file (`~/api_keys.sh`).
2. Run `claude-providers sync` to discover the key and create the alias.
3. `source ~/.local/share/claude-multi-account/aliases.sh` (or open a new shell).
4. Run `zai-coding-plan` to start a Claude Code session on `glm-5.2`.

No additional manual steps are needed — the key alias, model override, and endpoint are all defined in the provider config files.

Usage example

```
zai-coding-plan                                # launch Claude Code on glm-5.2
zai-coding-plan -p "explain this function"      # non-interactive
print mode
```

Xiaomi MiMo (xiaomi)

Xiaomi MiMo is Xiaomi's LLM platform. MiMo does expose a genuine **Anthropic-native endpoint** (<https://api.xiaomimimo.com/anthropic>, POST `/anthropic/v1/messages`), and the `xiaomi` alias originally used it — but as of **v1.19.0** the alias is pinned to **router transport** on MiMo's OpenAI-compatible `/v1` endpoint instead, along with every other hosted provider. Both paths were verified working; the uniform one is far easier to debug. Launches therefore go through `claude-code-router (ccr)`.

How the alias works

The key variable `XIAOMI_MIMO_API_KEY` in your keys file does **not** match the `models.dev xiaomi` provider's documented env name (`XIAOMI_API_KEY`), so `scripts/providers/key-aliases.json` maps `XIAOMI_MIMO_API_KEY` → `xiaomi`. An override in `scripts/providers/overrides.json` pins:

- **transport** `router`
- **base_url** `https://api.xiaomimimo.com/v1` (the OpenAI-compatible endpoint)
- **strong_model** `mimo-v2.5-pro`
- **fast_model** `mimo-v2.5`

The pinning is deliberate: `models.dev` lists a `mimo-v2.5-pro-ultraspeed` id that the **live API does not serve**, so the override guarantees only live-served model ids are used.

The resolved alias carries a **1048576-token** context (verify with `grep CMA_PROVIDER_CONTEXT_LIMIT ~/.local/share/claude-multi-account/providers/xiaomi.env`), which makes it one of the few aliases that comfortably absorbs a large enabled-plugin tool payload — see §9.2.

Models

MiMo's text-generation models (all support tool/function calling and reasoning; verified live 2026-06-19):

Model	Context	Reasoning	Tool Call	Notes
mimo-v2.5-pro	1M tokens	Yes	Yes	

Model	Context	Reasoning	Tool Call	Notes
				Flagship — strongest (alias strong)
mimo-v2.5	1M tokens	Yes	Yes	Omni / multimodal (image, audio, video) — alias fast
mimo-v2-pro	—	Yes	Yes	Legacy v2 Pro
mimo-v2-omni	256k tokens	Yes	Yes	Omni / multimodal (v2)
mimo-v2-flash	256k tokens	Yes	Yes	Cheapest tier (no longer the pinned fast model)

Non-text models exist (mimo-v2.5-asr, mimo-v2.5-tts*, mimo-v2-tts) but are speech-recognition / text-to-speech and cannot serve chat or code, so they are intentionally not wired as aliases.

Verified live

- GET /v1/models → HTTP 200, 10 models served (the 5 text models above + 5 ASR/TTS).
- POST /anthropic/v1/messages with Authorization: Bearer → HTTP 200, native Anthropic response (type: "message", content: [{type: "text"}, {type: "thinking"}]) for both mimo-v2.5-pro and mimo-v2-flash. (*The Anthropic-native endpoint works; the alias no longer uses it — see the transport note above.*)
- POST /v1/chat/completions with tools[] → finish_reason: "tool_calls", valid tool-call array, reasoning_content present (tool calling works).
- Streaming (stream: true) → SSE chat.completion.chunk deltas.
- Rate limits: 100 RPM / 10M TPM per account for text models (per-account, not per-key); expect 429 under load, retry with backoff.

Setup

Everything is already configured, but for a fresh install:

1. Ensure XIAOMI_MIMO_API_KEY is exported in your keys file (~/.api_keys.sh).
2. Run claude-providers sync to discover the key and create the alias.
3. source ~/.local/share/claude-multi-account/aliases.sh (or open a new shell).
4. Run xiaomi to start a Claude Code session on mimo-v2.5-pro.

Usage example

```
xiaomi                                # launch Claude Code on mimo-  
v2.5-pro  
xiaomi -p "explain this function" # non-interactive print mode
```

OpenCode Zen (provider id opencode, alias opencode-zen)

OpenCode Zen is OpenCode's curated AI gateway — a tested and verified list of models from multiple providers, accessed through a single API key. Zen includes **21 free models** (all \$0 cost, all support tool calling and reasoning), plus 49 paid models from OpenAI, Anthropic, Google, DeepSeek, and others.

The free models are available for a limited time while OpenCode collects feedback. One of them — **Big Pickle** — is a stealth model (its true identity is not disclosed; observed as deepseek-v4-flash behind the scenes).

How the alias works

The key variable `ZEN_API_KEY` (or `ApiKey_Opencode_Zen`) in your keys file is mapped to the opencode provider ID via `scripts/providers/key-aliases.json`. An override in `scripts/providers/overrides.json` pins:

- **alias** `opencode-zen` — so the launcher is `opencode-zen`, not the bare provider id. Hosts synced before that pin may still carry an opencode alias as well; `claude-providers` list is authoritative for what exists on yours.
- **strong_model** `big-pickle` — free stealth model, 200K context, reasoning + tool_call
- **fast_model** `deepseek-v4-flash-free` — free, 200K context, reasoning + tool_call

Those two are the *pins*. The per-model multi-alias phase (§5) additionally generates `opencode-zen1`, `opencode-zen2`, ... over the free-tier models it verifies, each with its own strong/fast pair, so the resolved model for a given alias is whatever its `.env` says — check with `claude-providers show <id>`.

Transport is **router** — the alias launches through `claude-code-router` (`ccr default-claude-code`), which translates the Anthropic protocol to the OpenAI-compatible Zen API (<https://opencode.ai/zen/v1/chat/completions>).

No transport or `base_url` override is needed — the `models.dev` catalog already has the correct values (`@ai-sdk/openai-compatible` → `router`, <https://opencode.ai/zen/v1>).

Free models available on Zen

All free models support tool calling and reasoning (verified live 2026-06-20):

Model	Context	Notes
big-pickle	200K	Stealth model — alias default (strong)
deepseek-v4-flash-free	200K	Alias fast model
mimo-v2.5-free	200K	Xiaomi MiMo
mimo-v2-pro-free	1M	Xiaomi MiMo Pro
mimo-v2-flash-free	262K	Xiaomi MiMo Flash
mimo-v2-omni-free	262K	Xiaomi MiMo Omni
nemotron-3-ultra-free	1M	NVIDIA Nemotron
nemotron-3-super-free	204K	NVIDIA Nemotron Super
north-mini-code-free	256K	North Mini Code
glm-4.7-free	204K	Zhipu GLM
glm-5-free	204K	Zhipu GLM 5
grok-code	256K	xAI Grok Code
kimi-k2.5-free	262K	Moonshot Kimi
minimax-m2.1-free	204K	MiniMax
minimax-m2.5-free	204K	MiniMax
minimax-m3-free	200K	MiniMax
qwen3.6-plus-free	262K	Alibaba Qwen
ring-2.6-1t-free	262K	Ring
hy3-preview-free	256K	HY3 Preview
ling-2.6-flash-free	262K	Ling (no reasoning)
trinity-large-preview-free	131K	Trinity (no reasoning)

Paid models start at \$0.05/1M tokens (GPT-5 Nano) and go up to \$30/1M (GPT-5.5 Pro). Claude models are available from \$1/1M (Haiku 4.5) to \$15/1M (Opus 4.1).

Verified live

- GET /v1/models → HTTP 200 (model list returned).
- POST /v1/chat/completions with big-pickle → HTTP 200, correct text response, cost=\$0, reasoning_content present. Stealth alias observed as deepseek-v4-flash.
- POST /v1/chat/completions with deepseek-v4-flash-free → HTTP 200, correct text, cost=\$0.

- Additional free models (mimo-v2.5-free, nemotron-3-ultra-free, north-mini-code-free) → all HTTP 200, all cost=\$0.
- Streaming (stream:true) → SSE chat.completion.chunk deltas.

Setup

Everything is already configured, but for a fresh install:

1. Ensure ZEN_API_KEY (or ApiKey_Opencode_Zen) is exported in your keys file (~/.api_keys.sh).
2. Run `claude-providers sync` to discover the key and create the alias.
3. `source ~/.local/share/claude-multi-account/aliases.sh` (or open a new shell).
4. Run `opencode-zen` to start a Claude Code session on big-pickle (via ccr).

Usage example

```
opencode-zen                                # launch Claude Code on
      big-pickle (via ccr)
opencode-zen -p "explain this function" # non-interactive print
      mode
claude-providers list                       # which opencode-zenN
      aliases exist here
```

Notes

- The alias uses **router transport** (ccr) because Zen's free models use OpenAI-compatible format, not Anthropic native format.
- Big Pickle is a stealth model — the actual model served may vary. This is by design per OpenCode's documentation.
- To use paid Claude models on Zen instead of free models, override `strong_model` and `fast_model` in `scripts/providers/overrides.json` (e.g. `claude-haiku-4.5` and `claude-sonnet-4-6`), and set transport to native with `base_url` to `https://opencode.ai/zen/v1/messages`.

Chutes (chutes)

Chutes is a pay-per-use AI inference platform with **TEE (Trusted Execution Environment)** models — all models run in secure enclaves for data privacy. The API is OpenAI-compatible at <https://llm.chutes.ai/v1>.

How the alias works

The key variable `CHUTES_API_KEY` in your keys file maps to the chutes provider ID. An override in `scripts/providers/overrides.json` pins:

- **strong_model** `zai-org/GLM-5.2-TEE` — newest GLM model, 202K context
- **fast_model** `Qwen/Qwen3.6-27B-TEE` — fast Qwen model, 262K context
- **context_limit** 262144 and **max_output** 65536

Transport is **router** — the alias launches through `claude-code-router` (`ccr default-claude-code`), which translates the Anthropic protocol to the OpenAI-compatible Chutes API.

Two things follow from those pins and are worth knowing before you debug a Chutes launch:

- Every Chutes model id is **vendor-prefixed** (`zai-org/...`, `Qwen/...`, `deepseek-ai/...`, `moonshotai/...`) — the exact shape that made 16 of 24 verified aliases unroutable before **v1.26.7**. If launches die as apparent timeouts, read §9.1 and confirm your router is current (`claude-ccr-build`).
- The **262144** context is the window §9.2 works through: on a host with a large enabled-plugin surface the tool payload alone can make it unusable regardless of the output cap.

Models available on Chutes

All models are TEE-enabled (Trusted Execution Environment):

Model	Context	Notes
zai-org/GLM-5.2-TEE	202K	Newest GLM — alias default (strong)
<code>zai-org/GLM-5.1-TEE</code>	202K	GLM 5.1
<code>zai-org/GLM-5-TEE</code>	202K	GLM 5
<code>Qwen/Qwen3.5-397B-A17B-TEE</code>	262K	Largest Qwen model
Qwen/Qwen3.6-27B-TEE	262K	Alias fast model
<code>Qwen/Qwen3-32B-TEE</code>	40K	Qwen 32B
<code>Qwen/Qwen3-235B-A22B-Thinking-2507-TEE</code>	262K	Thinking/reasoning model
<code>deepseek-ai/DeepSeek-V3.2-TEE</code>	131K	DeepSeek V3.2
<code>moonshotai/Kimi-K2.5-TEE</code>	262K	Kimi K2.5
<code>moonshotai/Kimi-K2.6-TEE</code>	262K	Kimi K2.6
<code>MiniMaxAI/MiniMax-M2.5-TEE</code>	196K	MiniMax M2.5

Model	Context	Notes
google/gemma-4-31B-turbo-TEE	131K	Google Gemma 4
unsloth/Mistral-Nemo-Instruct-2407-TEE	131K	Mistral Nemo

Important: Chutes is **pay-per-use** — the account must be funded to use models. Add balance at <https://chutes.ai>. The API key and configuration are correct, but requests will return "Quota exceeded" if the account balance is \$0.

Setup

1. Ensure CHUTES_API_KEY is exported in your keys file (~/.api_keys.sh).
2. Fund your Chutes account at <https://chutes.ai>.
3. Run `claude-providers sync` to discover the key and create the alias.
4. `source ~/.local/share/claude-multi-account/aliases.sh` (or open a new shell).
5. Run `chutes` to start a Claude Code session on `zai-org/GLM-5.2-TEE`.

Usage example

```
chutes                                # launch Claude Code on
    GLM-5.2-TEE (via ccr)
chutes -p "explain this function"    # non-interactive print mode
```

Verified

- API endpoint `https://llm.chutes.ai/v1/chat/completions` responds correctly
- Authentication with `Authorization: Bearer cpk_...` works
- All 13 TEE models are accessible (require funded account)
- OpenAI-compatible format confirmed

Poe (poe)

Poe is a universal AI platform with **382 models** from all major providers — chat, code, image generation, video generation, TTS, and STT. The API is OpenAI-compatible at <https://api.poe.com/v1>.

How the alias works

The key variable `POE_API_KEY` (or `ApiKey_Poe`) in your keys file maps to the poe provider ID. Transport is **router** — the alias launches through `claude-code-router` (`ccr default-claude-code`).

Aliases

Only one Poe alias is **pinned**; scripts/providers/overrides.json sets:

Alias	Strong Model	Fast Model	Focus
poe	claude-sonnet-4.6	gpt-5.4-mini	Primary — best Claude

Any additional poe2, poe3, ... aliases are **generated dynamically** by the per-model multi-alias phase (§5), not pinned here — and since that phase probes **free-tier models only** by default, a provider whose catalogue is entirely paid produces no extra aliases unless you opt in with --include-paid. Run `claude-providers` list to see what actually exists on your host rather than assuming the numbered aliases are present.

Model categories (382 total)

Category	Count	Examples
Chat/ Reasoning	130	claude-opus-4.8, gpt-5.5, grok-4, deepseek-v4-pro, gemini-3.1-pro, qwen3.7-max
Code	16	claude-code, gpt-5.3-codex, qwen3-coder-next, kimi-k2.7-code, seed-2.0-code
Image Generation	40	flux-2-pro, imagen-4, stable-diffusion3.5, qwen-image-2, grok-imagine-image
Video Generation	17	sora-2-pro, veo-3.1, kling-3.0, runway-gen-4.5, pixverse-v5.6
TTS/Voice	12	elevenlabs-v3, gemini-3.1-flash-tts, cartesia-ink-whisper, orpheus-tts
STT/Speech	1	whisper-v3-large-t
Other	166	perplexity-search, exa-research, hailuo, minimax-m3, glm-5.2, and more

Supported capabilities

- **Chat completions** — OpenAI-compatible /v1/chat/completions
- **Tool calling** — verified on claude-sonnet-4.6, gpt-5.4-mini, deepseek-v4-pro-e, grok-4
- **Streaming** — standard SSE streaming supported
- **Reasoning** — reasoning_content field in responses
- **Image input** — vision models support image URLs
- **Video generation** — via separate video endpoint

- **TTS/STT** — text-to-speech and speech-to-text models

Setup

1. Ensure POE_API_KEY is exported in your keys file (~/.api_keys.sh).
2. Run `claude-providers sync` to discover the key and create the alias.
3. `source ~/.local/share/claude-multi-account/aliases.sh` (or open a new shell).
4. Run `poe` to start a Claude Code session on `claude-sonnet-4.6` (via `ccr`).

Usage example

```
poe                                # launch Claude Code on
                                claude-sonnet-4.6 (via ccr)
poe -p "explain this function"    # non-interactive print mode
claude-providers list             # check which poeN aliases
                                exist on this host
```

Verified

- API endpoint `https://api.poe.com/v1/chat/completions` responds correctly
- Authentication with `Authorization: Bearer sk-poe-...` works
- Tool calling verified on multiple models
- 382 models accessible across all categories
- OpenAI-compatible format confirmed

HelixAgent (helixagent) — local single-GPU model

`helixagent` points Claude Code at a **local** HelixLLM backend — a podman container serving Qwen3-Coder-30B on one GPU (an RTX 5090) at `http://127.0.0.1:18434/v1` — instead of a hosted API.

Transport is **router**: launches go through `claude-code-router` plus the bundled Go `cma-proxy`, which recovers the model's Hermes/Qwen `<function=...>` tool calls into structured `tool_calls` so Claude Code's tools actually engage. The alias is pinned in `scripts/providers/helixagent.json` to a **229,376-token** context window and an **8,192-token** output cap.

Two things are specific to a local, single-GPU backend and worth knowing before you launch it.

The backend must be in "claude mode" (one big slot), not "coder mode"

The HelixLLM backend is shared — on the **same single GPU** — with HelixCode, and the two configurations are **mutually exclusive**:

Mode	llama.cpp config	Slot layout	For
claude mode	-c 229376 --parallel 1	one 229,376-token slot	the helixagent alias (Claude Code)
coder mode	-c 24576 --parallel 8	eight 3,072-token slots	HelixCode's concurrent sub-requests

llama.cpp splits -c across the parallel slots, so **coder mode gives each request only ~3,072 tokens** — a Claude Code session (system prompt + tool schemas) is far larger, and every launch against coder mode returns **HTTP 400**. helixagent needs the single large slot of claude mode.

Because coder mode is the **default** operational state (HelixCode is the common case), helixagent is honestly demoted to **unverified** and the launch gate refuses it until you flip the backend to claude mode:

```
# The mode-switch script lives in the companion helix_code repo –
  NOT this toolkit:
helix_code/scripts/helixllm-mode.sh claude
  # one 229,376-token slot
# then re-verify + launch:
claude-providers verify helixagent --deep      # passes once the
  backend is in claude mode
helixagent
```

Switch back with `helixllm-mode.sh coder` when you need HelixCode again. The two modes cannot run at once on the shared GPU.

Auto-start is opt-in — CMA_HELIX_AUTOSTART

Launching helixagent does **not** start HelixLLM for you by default. Opt in per launch (or export it) with any of the truthy values 1, true, yes, on — matched case-insensitively; anything else, unset included, means off:

```
CMA_HELIX_AUTOSTART=1 helixagent
```

The default is off because booting HelixLLM is a **machine-wide** side effect, not a private one: it claims the single GPU, and starting it in claude mode **evicts whatever HelixCode had running in coder mode**. A provider alias must not do that to a host merely because someone launched it.

Turning it off costs you no information. When HelixLLM is not answering, the pre-flight still says so and still prints the exact command:

```
helixagent: HelixLLM is not answering on http://127.0.0.1:18434
and auto-start is OFF by default.
```

Start it yourself (it claims the GPU and evicts HelixCode's coder mode):

```
helix_code/scripts/helixllm-mode.sh claude
```

Or opt in for this launch:

```
CMA_HELIX_AUTOSTART=1 <alias>
```

With the opt-in set, the wrapper searches the usual Helix Code locations (and \$HELIX_CODE_DIR) for `helixllm-mode.sh`, runs it in claude mode, then polls `/health` for up to 120 s — reporting a timeout with `podman logs helixllm-coder` and `<script> status` as next steps rather than proceeding silently. Auto-start only covers the *not running* case: a backend that is up but sitting in coder mode (`context_limit=24576`) gets the separate coder-mode warning above, and is never switched out from under HelixCode.

Minimal-launch (--bare) trim mode — CMA_PROVIDER_TRIM='bare'

A local 229,376-token window is small next to a hosted model's, and a normal launch would overflow it before the first turn: on a plugin-heavy host, auto-resumed session history (~330k tokens) plus the fixed hook/plugin/MCP/CLAUDE.md tool schemas (~110k tokens) together dwarf the whole window. So `helixagent`'s resolved env file carries `CMA_PROVIDER_TRIM='bare'`, which makes **every conversation launch minimal and fresh**:

- prepends **--bare**, dropping the hook/plugin/MCP/CLAUDE.md surface;
- **skips both automatic history seams** — the conversation-args `auto---` resume injection *and* the interactive (zero-args) stored session-flags injection — so no synced session history rides along;
- so each launch is a **fresh session that fits the local window**.

What trim deliberately does **not** touch:

- **Explicit** session selectors you pass yourself (`--resume <id>`, `--session-id <id>`, `--continue`, `--fork-session`, `-c`) are honored **verbatim** — trim only suppresses the *automatic* resume, never your own.

- **Non-conversation subcommands** (agents, mcp, export, doctor, config, plugin, setup, ...) get neither `--bare` nor a resume.
- **Untrimmed providers are byte-identical to before.** Trim is opt-in per provider via the `.env` line, wired today only for helixagent but reusable for any local-model provider that needs a fresh minimal session each launch.

13. Releasing — the mandatory live pre-release gate

claude-release-gate is the release gate: **no release commit may be made unless it exits 0**. It exists because a green sandbox suite is not proof the real host works — v1.25.1 shipped with the entire sandbox suite green while **every** router alias on the real host was bricked. The sandbox proves wrapper *logic*; it is structurally blind to real-host state (a PATH-shadowing ccr doppelgänger, a mis-configured local backend, resumed history overflowing a small window). So the gate adds a **live** layer that drives the real chain end-to-end.

```
claude-release-gate                # sandbox suite + live
    smoke (default provider: helixagent)
claude-release-gate --provider poe  # gate through a
    different provider
claude-release-gate --skip-suite    # reuse a suite run you
    JUST completed green
claude-release-gate --verify-providers # also run the full
    LLMsVerifier provider scan
```

The layers, fail-closed — any failure means **DO NOT RELEASE**:

1. **Sandbox suite** — `scripts/tests/run-all.sh`. `--skip-suite` skips it, only valid if you just ran it green.
2. **Live alias smoke** — regenerates the aliases from the *current* `lib.sh`, then launches the real alias with a fresh session (... `--session-id <uuid> -p "Reply with exactly: GATE-OK"`) through the real **PATH** → **ccr** → **route-apply** → **proxy** → **provider backend**, and asserts (a) the launch exited 0, (b) the served reply contained GATE-OK, and (c) for a router-transport provider, that the gateway's **sink-side route** (`.Router.default` in `~/.claude-code-router/<alias-id>/config.json`) actually names the provider under test. A write-then-apply route that silently served the wrong backend fails here.
3. **Provider scan** (*opt-in*, `--verify-providers`) — runs `claude-verify-providers` (the LLMsVerifier scan) over every provider/model. Slower; off by default.

The gate provider is chosen as `--provider`, else `$CMA_GATE_PROVIDER`, else `helixagent`. It **must exist and be verified** — a missing or broken gate provider is a gate **failure** (fix it, or pick another with `--provider`), never a silent skip. For `helixagent` specifically that means the HelixLLM backend must be in `claude` mode first (§12).